

CASCADEFORMER: A FAMILY OF TWO-STAGE CASCADING TRANSFORMERS FOR SKELETON-BASED HUMAN ACTION RECOGNITION

Anonymous authors

Paper under double-blind review

ABSTRACT

Skeleton-based human action recognition leverages sequences of human joint coordinates to identify actions performed in videos. Owing to the intrinsic spatiotemporal structure of skeleton data, Graph Convolutional Networks (GCNs) have been the dominant architecture in this field. However, recent advances in transformer models and masked pretraining frameworks open new avenues for representation learning. In this work, we propose **CascadeFormer**, a family of two-stage cascading transformers for skeleton-based human action recognition. Our framework consists of a masked pretraining stage to learn generalizable skeleton representations, followed by a cascading fine-tuning stage tailored for discriminative action classification. We evaluate CascadeFormer across three benchmark datasets—Penn Action, N-UCLA, and NTU RGB+D 60—achieving competitive performance on all tasks. To promote reproducibility, we will release our code and model checkpoints.

1 INTRODUCTION

Human action recognition is a fundamental computer vision task that aims to identify the action performed by a target person in a given video, represented as a sequence of frames. Human actions can be interpreted through various modalities, such as raw RGB images Zhang et al. (2021), skeletal joint coordinates Zhao et al. (2019), or a fusion of both Kim et al. (2023). Compared to RGB-based approaches, skeleton-based action recognition offers distinct advantages Zhou et al. (2023): it is computationally efficient due to relying solely on human joint coordinate data, and it is more robust to environmental noise and variations in camera viewpoints. Using these strengths, skeleton-based action recognition has been explored through various methods, including probabilistic models Zhao et al. (2019), recurrent neural networks (Shahroudy et al., 2016; Liu et al., 2016), and graph convolutional networks (GCN) Yan et al. (2018). Nevertheless, the rapid advancement of attention mechanisms Vaswani et al. (2023) in natural language processing has driven the widespread adoption of transformer architectures in computer vision. The success of vision transformers Dosovitskiy et al. (2021); Rao et al. (2021) across tasks such as image classification Ridnik et al. (2021), multimodal learning Radford et al. (2021), and visual instruction tuning Liu et al. (2023) underscores the potential of transformers for skeleton-based action recognition. Recent approaches incorporate novel attention designs tailored to spatiotemporal data Bertasius et al. (2021); Zhou et al. (2023); Wang et al. (2023); however, these transformer-based models are predominantly trained in an **end-to-end** manner.

However, instead of relying on end-to-end training, which may risk overfitting on downstream tasks, both language modeling Devlin et al. (2019) and multimodal learning Srivastava & Sharma (2023; 2024) have widely embraced masked pretraining as a precursor to supervised adaptation. **Motivated by the goal of equipping simple transformers with masked pretraining**, we propose **CascadeFormer**, a family of two-stage cascading transformers for skeleton-based human action recognition. CascadeFormer comprises a lightweight transformer for masked pretraining and an additional transformer for fine-tuning on action labels. In this paper, we begin with a review of prior work in masked pretraining with transformers and skeleton-based action recognition, followed by a brief overview of the input skeleton data format. We then present our key contributions:

- 054 1. We introduce **CascadeFormer**, a framework comprising three model variants that share a
 055 unified pipeline involving masked pretraining followed by cascading fine-tuning.
 056
 057 2. We conducted extensive evaluations of CascadeFormer on three widely used and diverse
 058 datasets: Penn Action, N-UCLA, and NTU RGB+D 60.
 059
 060 3. We performed extensive ablation studies, summarized in the main paper and further ex-
 061 panded in the Appendix, to analyze the impact of architectural choices and training config-
 062 urations.

062 Finally, we conclude with a discussion of the implications of our work.
 063

064 2 RELATED WORK 065

066 In this section, we first provide a brief overview of the development of masked pretraining frame-
 067 works in various fields. We then review widely adopted approaches in the recognition of human
 068 action based on skeletons.
 069

070 **Masked Pretraining Framework** The masked pre-training was first introduced in the domain of
 071 language modeling, most notably with BERT Devlin et al. (2019). By randomly masking 15% of
 072 the input tokens and training the model to predict them in a self-supervised manner, BERT can learn
 073 deep bidirectional representations Devlin et al. (2019). A similar paradigm has been extended to
 074 computer vision with masked autoencoders (MAEs) He et al. (2021), which apply a significantly
 075 higher masking ratio (up to 80%), enabling effective self-supervised pretraining for vanilla vision
 076 transformers Dosovitskiy et al. (2021). More recently, the masked pretraining framework has been
 077 rapidly adopted in multimodal learning. Models such as OmniVec Srivastava & Sharma (2023) and
 078 OmniVec2 Srivastava & Sharma (2024) take advantage of masked auto-encoding techniques to learn
 079 unified representations in diverse modalities, including image, text, video, and audio.

080 **Skeleton-based Action Recognition** Human action recognition can be explored through various
 081 modalities, including raw RGB images Zhang et al. (2021), skeletal joint coordinates Zhao et al.
 082 (2019), or a fusion of both Kim et al. (2023). In contrast to RGB-based methods, **skeleton**-based
 083 action recognition is characterized by its computational efficiency, owing to its reliance solely on
 084 human joint coordinate data, and its robustness to environmental noise and camera viewpoint vari-
 085 ations, since it preserves only the spatiotemporal information of keypoints Zhou et al. (2023). Al-
 086 though probabilistic models Zhao et al. (2019) demonstrated promising performance, deep learning
 087 approaches have become dominant in this domain. Given the temporal nature of video sequences,
 088 recurrent neural networks (RNN) Shahroudy et al. (2016); Liu et al. (2016) have proven effective in
 089 modeling temporal dynamics. Furthermore, convolutional neural networks (CNNs) Cai et al. (2023);
 090 Duan et al. (2022) have been adapted to this task by transforming skeleton data into pseudo-images,
 091 allowing convolutional operations to be applied. Furthermore, graph-based neural networks, includ-
 092 ing graph neural networks (GNNs) Shi et al. (2019) and graph convolutional networks (GCNs) Yan
 093 et al. (2018), have gained popularity due to the natural graph structure of human skeletons, where
 094 joints and bones are represented as vertices and edges, respectively. More recently, transformer ar-
 095 chitectures Vaswani et al. (2023) have attracted increasing attention in skeleton-based action recog-
 096 nition. These models are often enhanced with efficient spatiotemporal attention mechanisms Zhou
 097 et al. (2023); Do & Kim (2024); Bertasius et al. (2021); Wang et al. (2023), and are typically trained
 098 in an end-to-end fashion.

099 3 PRELIMINARIES 100

101 **Skeleton Data** Skeletons, or pose maps, refer to sets of Cartesian coordinates representing the
 102 key joints of the human body. Formally, given a batch of B video sequences, each consisting of
 103 T frames, with J joints per frame in a C -dimensional space, the input skeleton data X is typically
 104 structured as: $X \in \mathbb{R}^{B \times C \times T \times J}$. Although certain datasets such as Penn Action define skeletons in
 105 a two-dimensional space ($C = 2$), the most recent data sets—including N-UCLA and NTU RGB+D
 106 60—represent joint coordinates in a three-dimensional space ($C = 3$) to more accurately capture 3D
 107 motion. Since video lengths may vary across samples, it is common to either sample a fixed number
 108 of frames or apply dynamic padding based on the longest sequence in the batch.

108 **Multi-person Skeleton Data** Some datasets, such as NTU RGB+D 60, include actions involving
 109 multiple individuals. In these cases, assuming up to M persons appear in each video, the input
 110 skeleton data X can be extended to: $X \in \mathbb{R}^{B \times C \times T \times J \times M}$. When multiple persons are present,
 111 many approaches Do & Kim (2024); Zhou et al. (2023) encode all detected individuals. However,
 112 some methods Shahroudy et al. (2016) opt to retain only the most active person in the scene to
 113 reduce computational complexity or avoid noise from irrelevant actors.

114

115 4 CASCADEFORMER

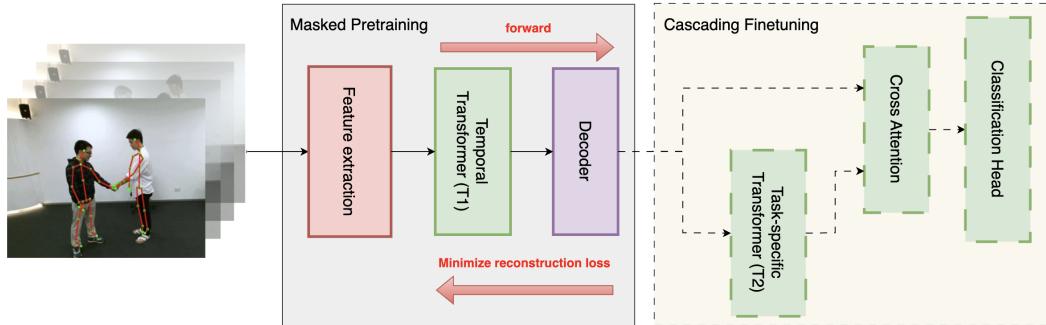
116

117 To equip vanilla transformers with masked pretraining, a technique widely adopted in language mod-
 118eling and multimodal learning, we propose **CascadeFormer**, a family of two-stage cascading trans-
 119 formers for skeleton-based human action recognition. As illustrated in Figure 1, CascadeFormer
 120 employs a lightweight transformer for masked pre-training, which is then followed by an additional
 121 transformer to fine-tune the action labels (Figure 2). We design a family of three model variants that
 122 differ in their feature extraction modules, as shown in Figure 3. The architecture is composed of
 123 three key components: masked pre-training, cascading fine-tuning, and feature extraction, each of
 124 which is described in detail in the following subsections.

125

126 4.1 MASKED PRETRAINING COMPONENT

127



128

129

130

131

132

133

134

135

136

137

138

139

140

Figure 1: **Overview of the masked pretraining component in CascadeFormer.** A fixed percentage of joints are randomly masked across all frames in each video. The partially masked skeleton sequence is passed through a feature extraction module to produce frame-level embeddings, which are then input into a temporal transformer (T1). A lightweight linear decoder is applied to reconstruct the masked joints, and the model is optimized using mean squared error over the masked positions. This stage enables the model to learn generalizable spatiotemporal representations prior to supervised finetuning.

141

142

143

144

145

146

147

148

The masked pretraining component, illustrated in Figure 1, enables the model to learn generalizable spatiotemporal dependencies in a self-supervised manner prior to any supervised training on action labels. Inspired by masked modeling strategies in other domains, we apply a joint-level masking scheme: specifically, 30% of the joints across all frames are randomly selected and masked by setting their coordinates to zero, while the remaining 70% retain their original Cartesian values. The partially masked skeleton sequence is then passed through a feature extraction module to generate frame-level embeddings. Although the three variants of CascadeFormer differ in their feature extraction designs, they all follow a common principle: each frame is treated as a token and projected into a high-dimensional embedding space. Formally, given a batch of B video sequences, each containing T frames with J joints per frame in a C -dimensional space, the resulting embedding tensor is denoted as $E \in \mathbb{R}^{B \times T \times \text{embed.dim}}$, where embed.dim denotes the transformer embedding dimension. These frame embeddings are then input into a vanilla temporal transformer backbone (T1), as shown in Figure 1. The output of T1, denoted as E_{pretrain} , maintains the same shape: $E_{\text{pretrain}} \in \mathbb{R}^{B \times T \times \text{embed.dim}}$. To reconstruct missing joint information, a lightweight linear decoder is applied to the output of T1. This decoder is used exclusively during pretraining and is

discarded during downstream fine-tuning. The model is trained to minimize the reconstruction error only in the masked joints. Let $\text{masked_}X \in \mathbb{R}^{B \times C \times T \times J}$ represent the masked input skeletons and $\text{masked_}X' \in \mathbb{R}^{B \times C \times T \times J}$ be the reconstructed output. The loss function is defined as the mean squared error (MSE) over the masked joints:

$$\mathcal{L}_{\text{MSE}} = \|\text{masked_}X - \text{masked_}X'\|^2$$

4.2 CASCADING FINETUNING COMPONENT

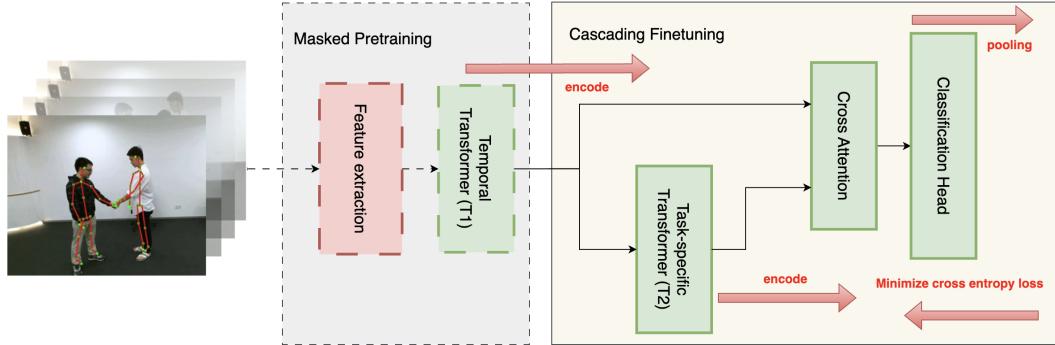


Figure 2: Overview of the cascading finetuning component in CascadeFormer. The frame embeddings produced by the pretrained temporal transformer backbone (T1) are passed into a task-specific transformer (T2) for hierarchical refinement. The output of T2 is fused with the original embeddings via a cross-attention module. The resulting fused representations are aggregated through frame-level average pooling and passed to a lightweight classification head. The entire model—including T1, T2, and the classification head—is optimized using cross-entropy loss on action labels during finetuning.

After masked pretraining, we introduce a **cascading** finetuning stage, as illustrated in Figure 2. This design is inspired by the hierarchical adaptation strategy proposed in OmniVec2 Srivastava & Sharma (2024) for multimodal learning. Unlike conventional fine-tuning, where a lightweight classification head is attached directly to a pre-trained transformer, CascadeFormer incorporates an additional task-specific transformer (T2) to refine features in a hierarchical manner.

Specifically, the frame embeddings $E \in \mathbb{R}^{B \times T \times \text{embed.dim}}$ obtained from the pretrained backbone (T1) are passed into the task-specific transformer (T2), which maps them into a new embedding space of the same dimensionality: $E_{\text{finetune}} \in \mathbb{R}^{B \times T \times \text{embed.dim}}$. These refined embeddings E_{finetune} are then fused with the original pretrained embeddings E_{pretrain} via a cross-attention module, as shown in Figure 2. Following the standard attention mechanism Vaswani et al. (2023), cross attention is computed as:

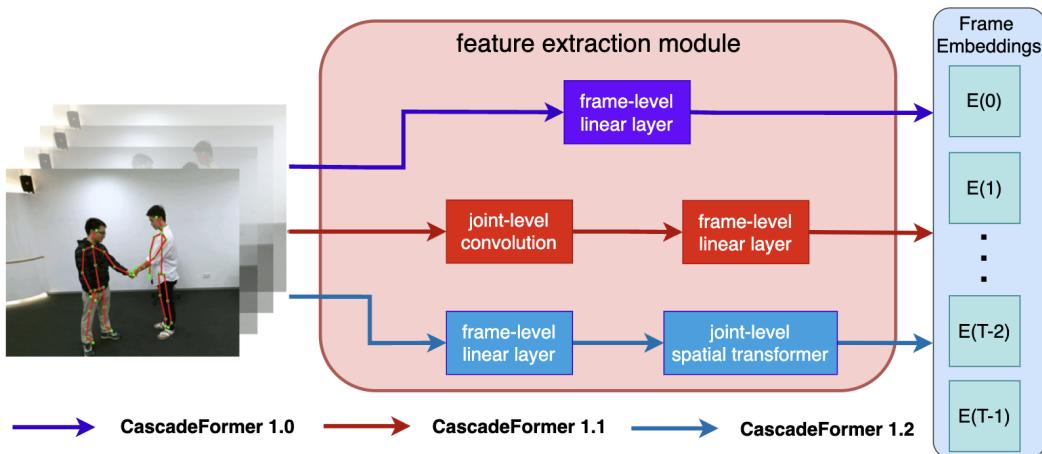
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$

$$E_{\text{cross}} = \text{Attention}(E_{\text{pretrain}}, E_{\text{finetune}}, E_{\text{finetune}})$$

To obtain a fixed-size video-level representation, we apply frame-level average pooling over the cross-attended embeddings:

$$E_{\text{avg}} = \frac{1}{T} \sum_{t=1}^T E_{\text{cross}}[:, t, :] \in \mathbb{R}^{B \times \text{embed.dim}}$$

Finally, a lightweight classification head is applied to the pooled embeddings, and the model is optimized using the standard cross-entropy loss with respect to the ground-truth action labels. All parameters of the pre-trained transformer backbone (T1), as well as the newly introduced T2 module and classification head, are trainable during the cascading fine-tuning stage.

216 4.3 FEATURE EXTRACTION MODULE
217

235 **Figure 3: Feature extraction module in CascadeFormer.** All variants convert input skeleton
236 sequences into frame-level embeddings for downstream temporal modeling. CascadeFormer 1.0 (purple)
237 applies a simple frame-level linear projection. CascadeFormer 1.1 (red) enhances this by first
238 applying a joint-level 1D convolution to capture spatial locality before linear projection. Cascade-
239 Former 1.2 (blue) constructs joint-level embeddings via a linear layer, refines them using a joint-level
240 spatial transformer, and aggregates the outputs into frame-level embeddings.
241

242 As illustrated in Figure 3, input skeleton sequences are first passed through a feature extraction mod-
243 ule to obtain preliminary frame-level embeddings E , which are then processed by the transformer
244 backbone ($T1$). We design a family of three CascadeFormer variants, each differing in how frame
245 embeddings are extracted:
246

247 **CascadeFormer 1.0** This baseline variant employs a simple linear projection to generate frame
248 embeddings directly from the input skeleton. Formally, given the input skeleton tensor $X \in \mathbb{R}^{B \times C \times T \times J}$, where B is the batch size, C the coordinate dimension, T the number of frames, and J
249 the number of joints, the linear projection outputs: $E \in \mathbb{R}^{B \times T \times \text{embed_dim}}$.
250

251 **CascadeFormer 1.1** To better capture spatial locality among joints, this variant prepends a
252 lightweight convolutional module prior to linear projection. The input skeletons $X \in \mathbb{R}^{B \times C \times T \times J}$
253 are reshaped in batch and temporal dimensions into $X \in \mathbb{R}^{(B \cdot T) \times C \times J}$, and then passed through a 1D
254 convolutional layer: $X_{\text{convolved}} \in \mathbb{R}^{(B \cdot T) \times C \times J}$. The resulting activations are passed into the same
255 linear projection as in CascadeFormer 1.0 to yield frame-level embeddings: $E \in \mathbb{R}^{B \times T \times \text{embed_dim}}$.
256

257 **CascadeFormer 1.2** In this variant, we first construct embeddings for each individual joint using a
258 linear layer, producing: $E_{\text{joint}} \in \mathbb{R}^{(B \cdot T) \times J \times (\text{embed_dim}/J)}$. A single layer spatial transformer (denoted
259 as ST) is then applied to the joint embeddings. $E_{\text{joint_ST}} \in \mathbb{R}^{(B \cdot T) \times J \times (\text{embed_dim}/J)}$. Finally, we
260 aggregate joint-level output into full frame-level embeddings $E \in \mathbb{R}^{B \times T \times \text{embed_dim}}$. Note that for
261 CascadeFormer 1.2, the total embedding dimension embed_dim must be divisible by the number of
262 joints J to allow for an even allocation across the joints. All three variants differ only in their feature
263 extraction mechanisms; the remainder of the architecture, including masked pre-training (Figure 1)
264 and cascading finetuning (Figure 2)—remains shared across all models.
265

266 5 EXPERIMENT SETUP
267

268 In this section, we introduce the datasets used for evaluation, followed by dataset-specific prepro-
269 cessing steps and training configurations.

270 5.1 DATASETS
271

272 We conduct a comprehensive evaluation of CascadeFormer on three widely-used benchmark datasets
 273 for human action recognition. **Penn Action** Zhang et al. (2013) dataset contains 2,326 video clips
 274 that span 15 human action classes. Each frame is annotated with 13 human joints in a 2D space.
 275 The dataset provides a standard 50/50 train/test split, ensuring mutually exclusive samples in both
 276 sets. **N-UCLA** wang et al. (2014) comprises 1,494 video clips categorized into 12 action classes.
 277 Skeletons are annotated with 20 joints in a 3D space. The dataset is collected from three camera
 278 viewpoints: two are used for training, and the third is reserved for evaluation, following the standard
 279 cross-view evaluation protocol. **NTU-RGB+D 60** Shahroudy et al. (2016), a large-scale dataset
 280 includes 56,880 video samples covering 60 action classes, with each skeleton containing 25 joints in
 281 3D space. NTU-RGB+D 60 offers two evaluation protocols: cross-subject (where subjects are split
 282 between training and testing) and cross-view (where specific camera angles are held out for testing).
 283 We adopted both protocols in our experiments.

284 5.2 DATA PREPROCESSING
285

286 Due to differences in scale, format, and split protocols across datasets, we apply dataset-specific
 287 preprocessing strategies. For **Penn Action**, we take advantage of the visibility flags provided with
 288 the annotations and remove occluded skeletons. Given the relatively small size of the dataset, we
 289 pad each video sequence within a batch to match the longest sequence in that batch. For **N-UCLA**,
 290 following the data pre-processing strategy in SkateFormer Do & Kim (2024), we virtually repeat
 291 the dataset multiple times to increase the diversity of the training. Each skeleton is normalized by
 292 centering it at the hip joint and applying scale normalization. Data augmentation techniques such
 293 as random rotation, scaling, and joint/axis dropping are applied. To avoid computational overhead
 294 from batch-level padding, we randomly sampled a fixed length of 64 frames per sequence. For
 295 **NTU-RGB+D 60**, we adopt the preprocessing pipeline from CTR-GCN Chen et al. (2021), which
 296 includes random rotation augmentation. To handle variable-length sequences, we uniformly sample
 297 64 frames around the temporal center of each clip.

298 5.3 TRAINING SETUP
299

300 All models are implemented in PyTorch Paszke et al. (2019) and trained on a single NVIDIA
 301 GeForce RTX 3090 GPU. We adopt dataset-specific hyperparameters to ensure optimal performance
 302 across different datasets. For **Penn Action** and **N-UCLA**, we use AdamW Loshchilov & Hutter
 303 (2019) as the optimizer. For masked pre-training, the learning rate is fixed at 1.0×10^{-4} . During the
 304 cascading fine-tuning stage, we reduce the learning rate to 1.0×10^{-5} and apply cosine annealing
 305 scheduling Loshchilov & Hutter (2017). For **NTU-RGB+D 60**, masked pretraining is performed
 306 with a constant learning rate of 1.0×10^{-4} and cosine annealing. During cascading fine-tuning, we
 307 switch to SGD Ruder (2017) as the optimizer. For cross-subject evaluation, we use a base learning
 308 rate of 3.0×10^{-5} , while for cross-view evaluation, we set it to 1.0×10^{-4} . The number of epochs
 309 for both pre-training and fine-tuning varies across the three CascadeFormer variants.

310 6 RESULTS
311312 6.1 DATASET-WISE EVALUATION
313

314 We present the performance of the CascadeFormer variants in three benchmark datasets in Table 1.
 315 In general, each variant demonstrates competitive performance, with certain variants better suited
 316 to specific datasets. In the Penn Action dataset, **CascadeFormer 1.0** achieves the highest accuracy
 317 at 94.66%, while both **1.1** and **1.2** closely follow at 94.10%. On the **N-UCLA** dataset, **Cascade-**
 318 **Former 1.1** leads with an accuracy of 91.16%, while variants **1.2** and **1.0** reach 90.73% and 89.66%,
 319 respectively. For the large-scale NTU RGB+D 60 dataset , which includes two evaluation protocols
 320 - cross-subject (CS) and cross-view (CV)—**CascadeFormer 1.0** again outperforms the others with
 321 **81.01%** (CS) and **88.17%** (CV). Variant 1.2 closely follows, while 1.1 lags slightly behind on both
 322 splits. These results suggest that, perhaps surprisingly, the simplest feature extraction design (Cas-
 323 cadeFormer 1.0 with a linear frame encoder) can be as effective, if not more, than more complex

designs involving convolutions or spatial transformers. All model checkpoints used in Table 1 are publicly released via HuggingFace for further analysis and reproducibility.

Model Variant	Penn Action	N-UCLA	NTU60/CS	NTU60/CV
CascadeFormer 1.0	94.66%	89.66%	81.01%	88.17%
CascadeFormer 1.1	94.10%	91.16%	79.62%	86.86%
CascadeFormer 1.2	94.10%	90.73%	80.48%	87.24%

Table 1: **Overall accuracy evaluation results of CascadeFormer variants on three datasets.** CascadeFormer 1.0 consistently achieves the highest accuracy on Penn Action and both NTU60 splits, while 1.1 excels on N-UCLA. All checkpoints are open-sourced for reproducibility.

6.2 MULTI-PERSON ACTION ANALYSIS

Unlike Penn Action and N-UCLA, which feature only single-person actions, NTU RGB+D 60 includes 11 classes involving interactions between two individuals (e.g., hugging, handshaking, pushing). Following the recommendation of the NTU authors Shahroudy et al. (2016), we retain the person with the greatest variance in motion for all actions. As shown in Table 2, **CascadeFormer 1.0** achieves 81.81% accuracy in two-person actions (CS split), slightly higher than the 80.82% obtained in single-person actions. This suggests strong generalization across subjects for interaction-based actions, despite a smaller number of classes. However, on the CV split, the model performs better on single-person actions (88.92%) than on two-person actions (84.86%), indicating that generalizing interactions across varying camera viewpoints remains more challenging than across different individuals. These findings reveal that cross-view generalization of social interactions may benefit from additional structural modeling of inter-person dynamics, which could be explored in future work.

Action Type	NTU60/CS accuracy	NTU60/CV accuracy
Single-person	80.82%	88.92%
Two-persons	81.81%	84.86%
Overall	<u>81.01%</u>	<u>88.17%</u>

Table 2: **Performance of CascadeFormer 1.0 on multi-person actions in NTU RGB+D 60.** Accuracy is broken down into single-person actions and two-person interactions under both cross-subject and cross-view splits.

7 PERFORMANCE COMPARISON

In this section, we compare the performance of the three CascadeFormer variants across different datasets with other representative models in skeleton-based action recognition.

Model	Penn Action accuracy
AOG	85.5%
HDM-BG	93.4%
CascadeFormer 1.0	94.66%
CascadeFormer 1.1	94.10%
CascadeFormer 1.2	94.10%

Table 3: **Comparison on Penn Action Zhang et al. (2013).** AOG Nie et al. (2015) is a hierarchical graph-based model. HDM-BG Zhao et al. (2019) is a probabilistic Bayesian approach. CascadeFormer variants achieve top accuracy, with 1.0 reaching the highest.

Table 3 compares the model accuracy in the Penn Action dataset Zhang et al. (2013). AOG Nie et al. (2015), a hierarchical graph model, reaches 85.5% accuracy. HDM-BG Zhao et al. (2019), a probabilistic model within a Bayesian framework, achieves 93.4%. All three CascadeFormer variants exceed 94% accuracy, with **CascadeFormer 1.0** achieving the best performance at **94.66%**.

These results demonstrate the effectiveness of CascadeFormer on compact 2D skeleton datasets like Penn Action. Table 4 shows the results on the N-UCLA dataset wang et al. (2014). ESV Liu et al. (2017), a CNN-based visualization method, obtains 86.09%, while Ensemble TS-LSTM Lee et al. (2017), an RNN-based model, reaches 89.22%. CascadeFormer 1.0 already exceeds this with 89.66%, and 1.1 and 1.2 further improve to **91.16%** and 90.73%, respectively. These results confirm that CascadeFormer performs robustly on small-scale 3D skeleton datasets as well.

Model	N-UCLA accuracy
ESV	86.09%
Ensemble TS-LSTM	89.22%
CascadeFormer 1.0	89.66%
CascadeFormer 1.1	91.16%
CascadeFormer 1.2	90.73%

Table 4: **Comparison on N-UCLA wang et al. (2014).** ESV Liu et al. (2017) is a CNN-based visualization method. TS-LSTM Lee et al. (2017) is a recurrent model ensemble. CascadeFormer 1.1 achieves the highest accuracy.

On the NTU RGB+D 60 dataset, we evaluate both cross-subject (CS) and cross-view (CV) splits. As shown in Table 5, ST-LSTM Liu et al. (2016) achieves 69.2% (CS) and 77.7% (CV), while ST-GCN Yan et al. (2018), a graph-based method, improves performance to 81.5% and 88.3%, respectively. CascadeFormer 1.0 delivers comparable performance—**81.01%** (CS) and **88.17%** (CV)—despite not using any explicit graph structure. CascadeFormer 1.2 follows closely behind, achieving 80.48% and 87.24%. These results validate CascadeFormer scalability and robustness in complex large-scale 3D datasets. Even without spatial graphs, our model competes with state-of-the-art graph convolutional approaches.

Model	N-UCLA accuracy	Model	NTU60/CS accuracy
ESV	86.09%	ST-LSTM	69.2%
Ensemble TS-LSTM	89.22%	ST-GCN	81.5%
CascadeFormer 1.0	89.66%	CascadeFormer 1.0	81.01%
CascadeFormer 1.1	91.16%	CascadeFormer 1.1	79.62%
CascadeFormer 1.2	90.73%	CascadeFormer 1.2	80.48%

Table 5: **Comparison on NTU RGB+D 60 (both Cross-Subject and Cross-View).** ST-LSTM Liu et al. (2016) uses recurrent modeling. ST-GCN Yan et al. (2018) adopts graph convolutions. CascadeFormer 1.0 achieves performance comparable to ST-GCN without graph structures.

Insights The empirical results across all three datasets provide a couple of insights into the design and performance of CascadeFormer: First, the effectiveness of the cascading fine-tuning strategy is consistent across model variants and datasets. CascadeFormer variants demonstrate clear improvements over strong baselines in both 2D and 3D skeleton-based action recognition tasks. Second, CascadeFormer demonstrates competitive scalability to large-scale and complex datasets. On NTU RGB+D 60, despite the absence of an explicit graph structure, CascadeFormer achieves competitive accuracy performance. This suggests that transformer-based architectures, when paired with effective pretraining strategies, can be comparable to graph-based methods in modeling human motion dynamics.

8 ABLATION HIGHLIGHTS

8.1 NECESSITY OF STRONG PRETRAINING

In this ablation study, we examine the necessity of employing a strong transformer backbone (T1) through masked pre-training. Specifically, we conducted a case study on **CascadeFormer 1.0** using the cross-subject split of the NTU RGB+D 60 dataset, a large-scale and complex benchmark that is well suited to reveal performance differences under varying pre-training strengths. We pre-train the T1 backbone for 1, 100, and 200 epochs, followed by 100 epochs of fine-tuning in all settings to

ensure fair comparison. A weak backbone pretrained for only 1 epoch yields an accuracy of 76.38%, while extending pretraining to 100 epochs results in a substantial improvement to 81.01%. Although training for 200 epochs yields a marginal further increase to 81.09%, the additional computational cost outweighs the negligible performance gain. Therefore, we adopted 100 epochs of masked pretraining for all subsequent experiments. This study confirms that strong backbone, achieved by sufficient masked pre-training, is necessary and beneficial for robust performance in challenging action recognition tasks.

# Epochs of Pretraining	NTU60/CS accuracy
1 epoch	76.38%
100 epochs	81.01%
200 epochs	81.09%

Table 6: **Effect of pretraining duration on NTU RGB+D 60 (cross-subject) performance.** We report accuracy on the cross-subject split after pretraining **CascadeFormer 1.0** for 1, 100, and 200 epochs. Performance improves significantly with longer pretraining, but plateaus after 100 epochs, indicating diminishing returns beyond this point.

8.2 PRETRAINING STRATEGY

In this ablation study, we further explore two alternative pretraining strategies using **CascadeFormer 1.0** on the Penn Action dataset Zhang et al. (2013), with results summarized in Table 7. Our proposed strategy applies random masking to 30% of joints across the entire video sequence, and optimizes a reconstruction loss. This method achieves a high accuracy of 94.66% on Penn Action. As a comparison, we evaluated a frame-level masking strategy, where 30% of entire frames are randomly masked. This results in a lower accuracy of 89.98%, indicating that reconstructing full frames is significantly more challenging for the transformer backbone than reconstructing scattered joint positions. We also investigate a **non-masked** variant in which the full skeleton sequence is reconstructed without any masking. This strategy achieves 91.10% accuracy, suggesting that while full reconstruction may be effective, it can lead to overfitting on smaller datasets, such as Penn Action. Based on these observations, we adopt the random joint masking strategy as our default pretraining approach for all subsequent experiments.

Pretraining Strategy	Penn Action accuracy
random joint masking	94.66%
random frame masking	89.98%
regular reconstruction	91.10%

Table 7: **Effect of pretraining strategy on Penn Action performance.** We compare three pretraining objectives for **CascadeFormer 1.0**: random joint masking, random frame masking, and regular (unmasked) reconstruction. Random joint masking yields the best downstream performance, highlighting the benefit of fine-grained spatial masking.

9 CONCLUSION

In this paper, we introduced CascadeFormer, a transformer-based model for skeleton-based action recognition with masked pretraining and cascading finetuning. Evaluated on Penn Action, N-UCLA, and NTU RGB+D 60, CascadeFormer consistently achieved strong performance on spatial modalities and dataset scales. Our results demonstrate the effectiveness of unified transformer pipelines for generalizable action recognition. Future directions include scaling to longer sequences, incorporating graph-aware modules, and extending to multimodal settings with RGB or depth inputs.

486 REFERENCES
487

- 488 Gedas Bertasius, Heng Wang, and Lorenzo Torresani. Is space-time attention all you need for video
489 understanding?, 2021. URL <https://arxiv.org/abs/2102.05095>.
- 490 Dongqi Cai, Yangyuxuan Kang, Anbang Yao, and Yurong Chen. Ske2grid: Skeleton-to-grid rep-
491 resentation learning for action recognition, 2023. URL <https://arxiv.org/abs/2308.07571>.
- 492
- 493 Yuxin Chen, Ziqi Zhang, Chunfeng Yuan, Bing Li, Ying Deng, and Weiming Hu. Channel-
494 wise topology refinement graph convolution for skeleton-based action recognition, 2021. URL
495 <https://arxiv.org/abs/2107.12213>.
- 496
- 497 Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep
498 bidirectional transformers for language understanding, 2019. URL <https://arxiv.org/abs/1810.04805>.
- 499
- 500 Jeonghyeok Do and Munchurl Kim. Skateformer: Skeletal-temporal transformer for human action
501 recognition, 2024. URL <https://arxiv.org/abs/2403.09508>.
- 502
- 503 Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas
504 Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit,
505 and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at
506 scale, 2021. URL <https://arxiv.org/abs/2010.11929>.
- 507
- 508 Haodong Duan, Yue Zhao, Kai Chen, Dahua Lin, and Bo Dai. Revisiting skeleton-based action
509 recognition, 2022. URL <https://arxiv.org/abs/2104.13586>.
- 510
- 511 Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked
512 autoencoders are scalable vision learners, 2021. URL <https://arxiv.org/abs/2111.06377>.
- 513
- 514 Sangwon Kim, Dasom Ahn, and Byoung Chul Ko. Cross-modal learning with 3d deformable atten-
515 tion for action recognition, 2023. URL <https://arxiv.org/abs/2212.05638>.
- 516
- 517 Inwoong Lee, Doyoung Kim, Seoungyoon Kang, and Sanghoon Lee. Ensemble deep learning for
518 skeleton-based action recognition using temporal sliding lstm networks. In *2017 IEEE Interna-*
519 *tional Conference on Computer Vision (ICCV)*, pp. 1012–1020, 2017. doi: 10.1109/ICCV.2017.
115.
- 520
- 521 Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning, 2023. URL
<https://arxiv.org/abs/2304.08485>.
- 522
- 523 Jun Liu, Amir Shahroudy, Dong Xu, and Gang Wang. Spatio-temporal lstm with trust gates for 3d
524 human action recognition, 2016. URL <https://arxiv.org/abs/1607.07043>.
- 525
- 526 Mengyuan Liu, Hong Liu, and Chen Chen. Enhanced skeleton visualization for view invariant
527 human action recognition. *Pattern Recognition*, 68:346–362, 2017. doi: 10.1016/j.patcog.2017.
02.030.
- 528
- 529 Ilya Loshchilov and Frank Hutter. Sgdr: Stochastic gradient descent with warm restarts, 2017. URL
<https://arxiv.org/abs/1608.03983>.
- 530
- 531 Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization, 2019. URL <https://arxiv.org/abs/1711.05101>.
- 532
- 533 Bruce Xiaohan Nie, Caiming Xiong, and Song-Chun Zhu. Joint action recognition and pose estima-
534 tion from video. In *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*,
535 pp. 1293–1301, 2015. doi: 10.1109/CVPR.2015.7298734.
- 536
- 537 Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan,
538 Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas
539 Kopf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy,
Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. Pytorch: An imperative style,

- 540 high-performance deep learning library. In *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 8024–8035, 2019. URL <https://papers.nips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html>.
- 541
- 542
- 543 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision, 2021. URL <https://arxiv.org/abs/2103.00020>.
- 544
- 545
- 546
- 547
- 548 Yongming Rao, Wenliang Zhao, Benlin Liu, Jiwen Lu, Jie Zhou, and Cho-Jui Hsieh. Dynamicvit: Efficient vision transformers with dynamic token sparsification, 2021. URL <https://arxiv.org/abs/2106.02034>.
- 549
- 550
- 551 Tal Ridnik, Emanuel Ben-Baruch, Asaf Noy, and Lihi Zelnik-Manor. Imagenet-21k pretraining for the masses, 2021. URL <https://arxiv.org/abs/2104.10972>.
- 552
- 553
- 554 Sebastian Ruder. An overview of gradient descent optimization algorithms, 2017. URL <https://arxiv.org/abs/1609.04747>.
- 555
- 556 Amir Shahroudy, Jun Liu, Tian-Tsong Ng, and Gang Wang. Ntu rgb+d: A large scale dataset for 3d human activity analysis, 2016. URL <https://arxiv.org/abs/1604.02808>.
- 557
- 558
- 559 Lei Shi, Yifan Zhang, Jian Cheng, and Hanqing Lu. Skeleton-Based Action Recognition with Di- rected Graph Neural Networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 7912–7921, Long Beach, CA, 2019. doi: 10.1109/CVPR.2019.00809.
- 560
- 561
- 562
- 563 Siddharth Srivastava and Gaurav Sharma. Omnivec: Learning robust representations with cross modal sharing, 2023. URL <https://arxiv.org/abs/2311.05709>.
- 564
- 565 Siddharth Srivastava and Gaurav Sharma. Omnivec2 - a novel transformer based network for large scale multimodal and multitask learning. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 27402–27414, 2024. doi: 10.1109/CVPR52733.2024.02588.
- 566
- 567
- 568
- 569 Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need, 2023. URL <https://arxiv.org/abs/1706.03762>.
- 570
- 571
- 572 Jiang wang, Xiaohan Nie, Yin Xia, Ying Wu, and Song-Chun Zhu. Cross-view action modeling, learning and recognition, 2014. URL <https://arxiv.org/abs/1405.2941>.
- 573
- 574
- 575 Qingtian Wang, Shuze Shi, Jabin He, Jianlin Peng, Tingxi Liu, and Renliang Weng. Iip-transformer: Intra-inter-part transformer for skeleton-based action recognition. In *2023 IEEE International Conference on Big Data (BigData)*, pp. 936–945. IEEE, December 2023. doi: 10.1109/bigdata59044.2023.10386970. URL <http://dx.doi.org/10.1109/BigData59044.2023.10386970>.
- 576
- 577
- 578
- 579
- 580 Sijie Yan, Yuanjun Xiong, and Dahua Lin. Spatial temporal graph convolutional networks for skeleton-based action recognition, 2018. URL <https://arxiv.org/abs/1801.07455>.
- 581
- 582 Chuhan Zhang, Ankush Gupta, and Andrew Zisserman. Temporal query networks for fine-grained video understanding, 2021. URL <https://arxiv.org/abs/2104.09496>.
- 583
- 584
- 585 Weiyu Zhang, Menglong Zhu, and Konstantinos G. Derpanis. From actemes to action: A strongly- supervised representation for detailed action understanding. In *2013 IEEE International Conference on Computer Vision*, pp. 2248–2255, 2013. doi: 10.1109/ICCV.2013.280.
- 586
- 587
- 588 Rui Zhao, Wanru Xu, Hui Su, and Qiang Ji. Bayesian hierarchical dynamic model for human action recognition. In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 7725–7734, 2019. doi: 10.1109/CVPR.2019.00792.
- 589
- 590
- 591 Yuxuan Zhou, Zhi-Qi Cheng, Chao Li, Yanwen Fang, Yifeng Geng, Xuansong Xie, and Margaret Keuper. Hypergraph transformer for skeleton-based action recognition, 2023. URL <https://arxiv.org/abs/2211.09590>.
- 592
- 593

594 **A APPENDIX**

595

596 We conduct further ablation studies to examine the effects of input representations, decoder archi-
 597 tectures, and backbone freezing strategies.

598

599 **A.1 INPUT DATA REPRESENTATION**

600

601 We investigate alternative input data representations beyond the original joint coordinates using
 602 **CascadeFormer 1.0**, as shown in Table 8. Surprisingly, the use of raw joint coordinates yields the
 603 highest accuracy of **94.66%** on the Penn Action dataset. To explore the effectiveness of bone-based
 604 representations, we develop three variants. The first approach constructs each bone by subtracting
 605 the coordinates of one joint from its adjacent joint, resulting in an accuracy of 92.32%. The second
 606 method concatenates the coordinates of two consecutive joints, achieving 93.16% accuracy. The
 607 third approach linearly parameterizes each bone segment using its slope and intercept, which attains
 608 93.91% accuracy. Overall, although all bone-based variants perform competitively, the original joint
 609 representation proves to be the most effective for our model on this task.

610

Data Representation	Accuracy
Joints	94.66%
Bones (Subtraction)	92.32%
Bones (Concatenation)	93.16%
Bones (Parameterization)	93.91%

611 Table 8: **Comparison of different input data representations on Penn Action.** Using the original
 612 joint coordinates achieves the highest accuracy (94.66%), outperforming all three alternative bone-
 613 based variants.

614

615 **A.2 DECODER ARCHITECTURE**

616

617 We conduct an ablation study to compare alternative decoder architectures for masked pretraining on
 618 the *CascadeFormer 1.0* using the Penn Action dataset. As shown in Table 9, our default choice—a
 619 simple linear layer to reconstruct masked joints—achieves the highest accuracy of **94.66%**. We then
 620 evaluate an MLP decoder composed of two linear layers with a ReLU activation in between, which
 621 yields a reduced accuracy of 92.51%. Finally, we test an MLP decoder with a residual connection
 622 to facilitate gradient flow, resulting in an even lower accuracy of 91.20%. These results suggest that
 623 the linear decoder not only provides the most effective reconstruction but also generalizes better,
 624 likely due to its lower risk of overfitting on small-scale datasets. Based on this finding, we adopt the
 625 linear decoder for all experiments throughout this work.

626

Decoder Architecture	Accuracy
linear	94.66%
MLP	92.51%
MLP + residual	91.20%

627 Table 9: **Comparison of decoder architectures during masked pretraining on Penn Action.** A
 628 simple linear decoder outperforms both MLP and MLP with residual connection, indicating that
 629 increased decoder complexity may lead to overfitting on smaller datasets.

630

631 **A.3 BACKBONE FREEZING DECISION**

632

633 In this ablation study, we examine the impact of parameter-freezing strategies for the transformer
 634 backbone during the cascading finetuning stage. Using *CascadeFormer 1.0* on the Penn Action
 635 dataset, we present our findings in Table 10. Our primary approach involves fully finetuning the
 636 entire backbone, allowing all transformer parameters to be updated during training. This strategy
 637 yields the best performance with an accuracy of **94.66%**. In contrast, freezing the entire backbone
 638 results in a significant drop in accuracy to 85.11%. We also explore partial finetuning by training
 639 only the final transformer layer, which achieves 88.39% accuracy. These results suggest that full
 640

648 backbone finetuning is crucial for effective downstream adaptation in action recognition. Conse-
 649 quently, we adopt full backbone finetuning for all experiments.
 650

Backbone Freezing Decision	Accuracy
fully finetune	94.66%
fully freeze	85.11%
finetune the last layer	88.39%

651
 652 **Table 10: Effect of different backbone freezing strategies during cascading finetuning on Penn**
 653 **Action.** Fully finetuning the transformer backbone yields the highest accuracy, while freezing all
 654 layers significantly degrades performance.
 655

656
 657
 658
 659
 660
 661
 662
 663
 664
 665
 666
 667
 668
 669
 670
 671
 672
 673
 674
 675
 676
 677
 678
 679
 680
 681
 682
 683
 684
 685
 686
 687
 688
 689
 690
 691
 692
 693
 694
 695
 696
 697
 698
 699
 700
 701